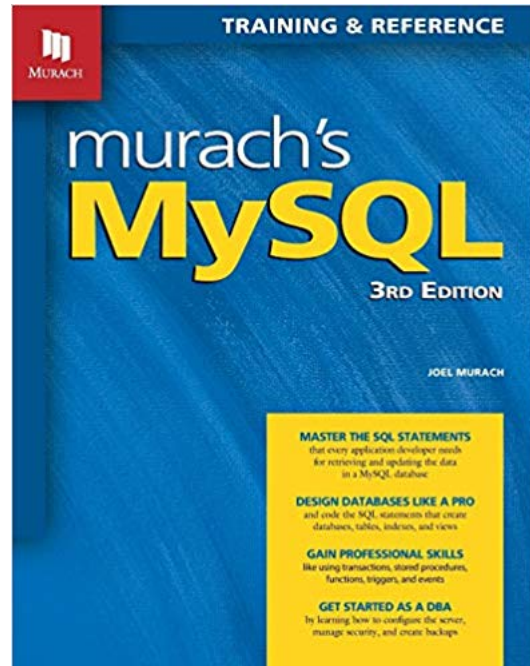

Subqueries or Nested queries

Topic 3

Lesson 5 – Building up a query

Chapter 7 Murach's MySQL



Getting the results, you want

Sometimes, it takes many transformations of the data in a database to get the results you want.

Nesting queries allows you to develop each of the transformations separately.

This allows you to take a divide and conquer approach to creating your queries.

This approach is a more modular approach, leading to a simpler approach for both building and testing your queries.

Example of a nested query

Suppose you want to find all the students that have earned above the average number of earned credits.

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

Step 1: Find the average number of earned credits for students

Step 2: Find all students who have earned more than the average calculated in step 1.

Solution: nested query

Step 1: Find the average number of credits

```
SELECT AVG(credits_earned) FROM student;
```

Step 2: Find all students who have earned more than the average calculated in step 1.

```
SELECT sid, name, credits_earned FROM student  
WHERE credits_earned >  
(SELECT AVG(credits_earned) FROM student);
```

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

A subquery can appear ...

1. In a WHERE clause as a search condition
 2. In a HAVING clause as a search condition
 3. In the FROM clause as a table specification
 4. In the SELECT clause as a column specification
- These are the only clauses where a subquery can appear.
 - A subquery can return a single value, a list of values or a table of values.
 - Subqueries are surrounded by parentheses.
 - A subquery may contain other subqueries.
 - Our previous example had a subquery in the WHERE clause.

Subquery Example in the WHERE clause

A subquery can be used to collect the values that you want to match for a criteria from a table.

Suppose you want to find all the student ids who are majoring in CS or DS.

student

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

Available_major

MID	MAJOR
1	CS
2	DS
3	Accounting

student_major

SID	MID
1	1
1	3
2	1
3	2

Solution: A nested query with IN

EXAMPLE: SELECT sid FROM student_major
WHERE mid IN
(SELECT mid from available_major WHERE
major = 'CS' or major = 'DS')

available_major

MID	MAJOR
1	CS
2	DS
3	Accounting

student_major

SID	MID
1	1
1	3
2	1
3	2

SID
1
2
3

What about solving this with a JOIN?

Solution: with a JOIN

```
EXAMPLE: SELECT sid FROM student_major
         JOIN available_major USING (mid)
         WHERE major = 'CS' OR major = 'DS' ;
```

available_major

MID	MAJOR
1	CS
2	DS
3	Accounting

student_major

sid	mid
1	1
1	3
2	1
3	2

result

SID
1
2
3

IF all result fields for a JOIN are from one table this is known as a **SEMIJOIN**. It will generate a simpler query plan than a typical join.

ANY and ALL keywords

Many times a subquery will return more than one value, however the tuple only has one value. You can address this discrepancy by using the keywords: SOME, ANY, or ALL.

ANY - compares a value to each value in the result list from a subquery and evaluates to TRUE if the conditional is true for ANY of the values. SOME provides the same functionality.

ALL - compares a value to each value in the result list from a subquery and evaluates to TRUE if the conditional is true for ALL of the values

Example with ALL keyword in WHERE clause

Write a query to find student ids that have credits_earned greater than all values for credits_earned for all CS majors.

student

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

Available_major

MID	MAJOR
1	CS
2	DS
3	Accounting

student_major

sid	mid
1	1
1	3
2	1
3	2

Solution: ALL keyword

```
SELECT sid from student WHERE credits_earned > ALL
(SELECT credits_earned FROM student WHERE sid IN
(SELECT sm.sid FROM available_major a JOIN
student_major sm USING (mid) WHERE major = 'CS' );
```

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

student

student_major

Available_major

MID	MAJOR
1	CS
2	DS
3	Accounting

solution

id

sid	mid
1	1
1	3
2	1
3	2

Example: ANY keyword

If we wanted to find the students that had credits_earned greater than any of the CS majors then use the ANY keyword.

```
SELECT sid from student
    WHERE credits_earned > ANY
        (SELECT credits_earned FROM student
        WHERE sid IN
            (SELECT sm.sid FROM available_major a
            JOIN student_major sm
            USING (mid) WHERE major = 'CS' );
```

solution

id

3

Uncorrelated and correlated queries

In all examples, our subquery is computed in one stage and our outer query is computed in a following stage. This is known as an uncorrelated query. The inner query is computed once, and its result values are used by the outer query which is also computed once.

Sometimes you may want the inner and outer queries to be computed in the same phase of the query. This is known as a correlated query.

With a correlated query the inner query has access to the field values in the outer query. For each tuple in the outer query the inner query is computed. A correlated query is comparable to executing a loop in a host language.

Correlated query example

Find students who have earned `credits_earned` greater than the average `credit_earned` for each student's respective major.

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020

sid	mid
1	1
1	3
2	1
3	2

mid	AVG(credits_earned)
1	48
2	50
3	32

Correlated query solution

Find students and their major who have earned credits_earned greater than the average credit_earned for each student's respective major.

```
SELECT s.sid, sm.mid FROM student AS s JOIN
      student_major AS sm USING (sid)
      WHERE credits_earned >
            (SELECT AVG(credits_earned)
              FROM student s2 JOIN
                    student_major AS sm2
                  ON s2.id = sm2.id
              WHERE sm.mid = sm2.mid );
```


EXISTS keyword

Sometimes you are not interested in the actual values returned from a subquery, you are interested if there are any rows in the table that satisfies the condition of the subquery. If so, use the EXISTS keyword.

Syntax:

WHERE [NOT] EXISTS (subquery)

If the subquery returns a result, the conditional result is TRUE, if the subquery does not return a result then the result of the WHERE clause is FALSE.

WHERE EXISTS tests if one or more rows are returned by the subquery, if results, returns TRUE if not FALSE.

WHERE NOT EXISTS tests that no rows are returned by the subquery, if no result then returns TRUE if not FALSE .

Example EXISTS keyword

Write a query to find student ids that have not declared a major

student					
sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020
4	Frosh	COS	8	128	2025

Available_major	
mid	major
1	CS
2	DS
3	Accounting

student_major	
sid	mid
1	1
1	3
2	1
3	2

Solution: NOT EXISTS

```
SELECT sid FROM student AS s WHERE NOT EXISTS  
(SELECT sid from student_major AS sm  
WHERE sm.sid = s.sid);
```

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020
4	Frosh	COS	8	128	2025

sid	mid
1	1
1	3
2	1
3	2

sid
4

Exercise: Solve this with a JOIN

Write a query to find student ids that have not declared a major

student					
sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020
4	Frosh	COS	8	128	2025

Available_major	
mid	MAJOR
1	CS
2	DS
3	Accounting

student_major	
sid	mid
1	1
1	3
2	1
3	2

Solution: LEFT OUTER JOIN

```
SELECT s.sid, sm.sid FROM student s
LEFT JOIN student_major sm
ON s.sid = sm.sid WHERE sm.sid IS NULL;
```

sid	Name	School	Credits_Earned	Credits_Req	Yr_grad
1	Smith	Khoury	32	120	2019
2	Shah	D'Amore McKim	64	128	2019
3	Li	Khoury	50	120	2020
4	Frosh	COS	8	128	2025

sid	mid
1	1
1	3
2	1
3	2

s.sid	sm.sid
4	NULL

A Subquery in the SELECT clause

Many times you want to report an aggregated value for each tuple in a table. Use a subquery in the SELECT clause to create this type of result.

If a subquery is in the SELECT clause it must return a single value.

Subqueries in the SELECT clause are typically correlated.

EXAMPLE: Find the number of majors for each student, report the student id, student name and number of majors for each student.

EXAMPLE: Subquery in the SELECT clause

EXAMPLE: Find the number of majors for each student, report the student id, student name and number of majors for each student.

Use the student_major table to find the COUNT of majors for each student

SOLUTION:

```
SELECT sid, name, ( SELECT COUNT(*) FROM
student_major AS sm WHERE sid = s.sid )
AS num_majors FROM student AS s;
```

A Subquery in the FROM clause

An inline view is a SELECT statement in the FROM clause.

An in-line view is commonly used to simplify complex queries by removing join operations and condensing several separate queries into a single query.

All subqueries in the FROM clause must be named or given an alias. You should also alias (provide a name) for any field that is a computed

This approach allows queries to evolve into a solution, making it easier to test and validate the subqueries.

EXAMPLE: A Subquery in the FROM clause

Find the majors for each student, report the student id, student name and major names for each student.

```
SELECT s.sid, name, major FROM student AS s
JOIN
  ( SELECT sid, am.mid, major
    FROM student_major AS sm
    JOIN available_majors AS am
      ON sm.mid = am.mid ) AS mm
ON s.sid = mm.sid;
```

Result: A Subquery in the FROM clause

```
SELECT s.sid, name, major FROM student AS s
JOIN
( SELECT sid, am,mid, major
  FROM student_major AS sm
  JOIN available_majors AS am
    ON sm.mid = am.mid ) AS mm ON s.sid = mm.sid;
```

sid	Name	major
1	Smith	CS
1	Smith	ACCOUNTING
2	Shah	CS
3	Li	DS

This provides one row per student major. What if you wanted one row per student and the Major values combined ?

Result: A tuple per student

```
SELECT s.sid, name, GROUP_CONCAT(major)
FROM student AS s
JOIN
( SELECT sid, am, mid, major
  FROM student_major AS sm
  JOIN available_majors AS am
    ON sm.mid = am.mid ) AS mm
ON s.sid = mm.sid GROUP BY id, name;
```

sid	Name	major
1	Smith	CS,Accounting
2	Shah	CS
3	Li	DS

What if we wanted to report on all students?

Result: A tuple for every student

```
SELECT s.sid, name, GROUP_CONCAT(major)
FROM student AS s
LEFT JOIN
( SELECT sid, am.mid, major
  FROM student_major AS sm
  JOIN available_majors AS am
    ON sm.mid = am.mid ) AS mm
ON s.sid = mm.sid GROUP BY id, name;
```

sid	Name	major
1	Smith	CS,Accounting
2	Shah	CS
3	Li	DS
4	Frosh	NULL

Advantages of different approaches

SUBQUERY

Allows you to pass a value to the main query

Queries involving many tables may be easier to comprehend

Easier to comprehend when the query is not focused on the foreign key/primary key relationship

May lead to faster processing time

JOIN

Access to all fields in all tables at the same level

Easier to comprehend when the query is focused on the foreign key/primary key relationship

Practice work:

Let's download the subqueries exercises from canvas and work in groups to complete them.

Advanced topic

Common Table Expressions

Common Table Expression (CTE)

Nesting queries within the FROM clauses can lead to an SQL query that is difficult to read.

The WITH construct was introduced in ANSI SQL Standard 1999 and is supported in MySQL 8.x. It allows you to define and name a temporary result that can be used within the following SQL statement.

Common Table Expression: Syntax

With a CTE, all your IN-line tables are defined at the beginning of your query. Your **sql_statement** can treat your CTE tables as base tables.

SYNTAX:

```
WITH [RECURSIVE] cte_name1  
  AS (subquery1)  
  [, cte_name2 AS (subquery2)]  
  [...]  
  sql_statement
```

EXAMPLE: No CTE vs. CTE

```
SELECT id, name, major FROM student AS s
JOIN
( SELECT sid, mid, major
  FROM student_major AS sm
  JOIN available_majors AS am
    ON sm.mid = am.mid ) AS mm
ON s.sid = mm.sid;
```



No
CTE



CTE

```
WITH mm AS
( SELECT sid, mid, major FROM student_major AS sm
  JOIN available_majors AS am
    ON sm.mid = am.mid )
SELECT id, name, major FROM student AS s
JOIN mm ON s.sid = mm.sid;
```

Recursive CTEs

Recursive CTEs allows a query to process a hierarchical relationship represented in the data

A recursive CTE has a base case and a recursive case. The result of the recursive CTE is the UNION of the base case and the recursive case.

The query will continue to run until there are no tuples left to process.

Example: Recursive CTE

Let's say you want to trace through the manager – subordinate relationship in the employee table. It is a hierarchal relationship, where the number of levels in the organization is not limited. In order to do this we need some mechanism that will continue to read the tuples until we have reached the employees who have no one reporting to them.

	employee_id	last_name	first_name	department_number	manager_id
▶	1	Smith	Cindy	2	NULL
	2	Jones	Elmer	4	1
	3	Simonian	Ralph	2	2
	4	Hernandez	Olivia	1	9
	5	Aaronsen	Robert	2	4
	6	Watson	Denise	6	8
	7	Hardy	Thomas	5	2
	8	O'Leary	Rhea	4	9
	9	Locario	Paulo	6	1

Walk the hierarchy

	employee_id	last_name	first_name	manager_id
►	1	Smith	Cindy	NULL
	2	Jones	Elmer	1
	3	Simonian	Ralph	2
	4	Hernandez	Olivia	9
	5	Aaronsen	Robert	4
	6	Watson	Denise	8
	7	Hardy	Thomas	2
	8	O'Leary	Rhea	9
	9	Locario	Paulo	1

No boss = Cindy Smith
(1)

Base case
(1)

Reports to Cindy (1):

Elmer (2) Paolo (9)

First recursive
case (2)

Reports to :
Elmer (2) Paolo (9)

Ralph (3) Rhea (8)
Thomas (7) Olivia (4)

Second recursive
case (3)

Reports to :
Rhea (8) Olivia (4)

Denise (6) Robert (5)

Third recursive
case (4)

Recursive RTE to walk the hierarchy

```
WITH RECURSIVE employees_cte AS
(
    -- Nonrecursive query
    SELECT employee_id,
           CONCAT(first_name, ' ', last_name) AS employee_name,
           1 AS ranking
    FROM employees
    WHERE manager_id IS NULL
    UNION ALL
    -- Recursive query
    SELECT employees.employee_id,
           CONCAT(first_name, ' ', last_name),
           ranking + 1
    FROM employees
        JOIN employees_cte
        ON employees.manager_id = employees_cte.employee_id
)
SELECT *
FROM employees_cte
ORDER BY ranking, employee_id
```

Recursive CTE result

	employee_id	employee_name	ranking
▶	1	Cindy Smith	1
	2	Elmer Jones	2
	9	Paulo Locario	2
	3	Ralph Simonian	3
	4	Olivia Hernandez	3
	7	Thomas Hardy	3
	8	Rhea O'Leary	3
	5	Robert Aaronson	4
	6	Denise Watson	4

Summary

In this module you learned:

- Subqueries can be found in the SELECT, FROM, WHERE and HAVING clauses
- Correlated queries
- EXISTS
- NOT EXISTS
- ANY and ALL keywords